

## **Technical Assignment: Build a Mini LLM-Powered Question-Answering System Using RAG**

### **Objective**

You are tasked with building a simple, functional prototype of a document-based Question-Answering (QA) system using Retrieval-Augmented Generation (RAG) powered by an open-source LLM.

This is a time-boxed challenge (4 hours) designed to evaluate your ability to:

- Think like a builder under pressure
  - Design a clean RAG pipeline
  - Integrate embeddings, vector search, and LLMs
  - Use AI tools like ChatGPT or Copilot judiciously
- 

### **Provided Resources**

- Input Dataset: A text document or multiple documents will be provided by Wundrsight. These may include excerpts from clinical guidelines, mental health protocols, or structured medical text.
  - No need to search or scrape external data sources.
- 

### **Assignment Scope**

#### **Required Components**

Please implement the following in your solution:

1. Document Ingestion and Chunking
  - Load the provided text(s)
  - Split into manageable chunks for embedding (e.g., 200–500 tokens)
  - Clearly explain the chunking strategy in your code or documentation

## 2. Embedding and Vector Store

- Generate embeddings using a pre-trained model (e.g., all-MiniLM-L6-v2 from Sentence Transformers)
- Store vectors using FAISS or ChromaDB
- Ensure retrieval of top-k similar chunks for a given query

## 3. Query Interface

- Accept a text query from the user (e.g., via command line or notebook input)
- Retrieve relevant context chunks using semantic search

## 4. LLM Integration

- Use a local or free-tier accessible open-source LLM (e.g., LLaMA2, GPT4All, Mistral-7B, or any suitable HuggingFace-hosted model)
- Use the retrieved context to generate an answer to the input query
- You may use LangChain or other wrappers if needed, but explain your approach

## 5. Output

- Display or print the generated answer
- Include a sample input and output using this test question:

*Give me the correct coded classification for the following diagnosis:  
"Recurrent depressive disorder, currently in remission"*

---

### Time Limit

- You must complete the coding portion within 4 hours
- Use the README to note what is complete, what is skipped, and why
- You may use ChatGPT or Copilot to assist—but clearly call out where and how you used them

---

## Deliverables

Please submit:

1. Code (in Jupyter Notebook or .py file) with clear structure and comments
  2. README or explanatory note with:
    - Tools and models used
    - Where AI tools like ChatGPT/Copilot were used
    - Your design decisions and any assumptions made
    - Limitations due to time or resource constraints
  3. Sample output for the provided query
- 

## Evaluation Criteria

Area	Evaluation Focus
Functionality	End-to-end working pipeline (document → retrieval → LLM → answer)
Technical Design	Soundness of architectural decisions and justification
Use of AI Tools	Responsible and smart use of ChatGPT or Copilot
Code Quality	Modular, readable, and logically organized

Time Management

How well you prioritized features within time limit

Clarity

Assumptions and design decisions are clearly documented

---

**Optional Bonus (if time permits)**

- Add a simple Gradio or Streamlit interface
- Allow for dynamic document uploads
- Implement basic caching for repeat queries
- Add relevance reranking logic (e.g., using Maximal Marginal Relevance)

---

**Important Notes**

- If something is unclear, assume a reasonable approach and note it in the README.
- Do not overengineer. Focus on delivering a working, clean MVP.
- Do not exceed the 4-hour coding time limit—document anything left incomplete.